

Neural Network Basics

Binary Classification

Case: Input an image, Output a label.

Notation

One training example: (x, y) , $x \in \mathbb{R}^{n_x}$, $y \in \{0, 1\}$.

m training examples: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$.

$$X = \begin{bmatrix} | & | & & | \\ x^{(1)} & x^{(2)} & \dots & x^{(m)} \\ | & | & & | \end{bmatrix}$$

$$X \in \mathbb{R}^{n_x \times m}.$$

$$Y = [y^{(1)}, y^{(2)}, \dots, y^{(m)}]$$

$$Y \in \mathbb{R}^{1 \times m}$$

call `X.shape` and `Y.shape` in python.

Logistic Regression

Given x , want $\hat{y} = P(y = 1|x)$.

$$X \in \mathbb{R}^{n_x}, 0 \leq y \leq 1.$$

Parameters: $w \in \mathbb{R}^{n_x}$, $b \in \mathbb{R}$. For the output $w^T X + b$.

In order to constrain \hat{y} between 0 and 1, we need a sigmoid function $\sigma(w^T X + b)$.

$$\text{Sigmoid Function: } \sigma(z) = \frac{1}{1 + e^{-z}}.$$

The goal of logistic regression is **trying to learn parameters w and b so that \hat{y} has become a good estimate of the chance of y being equal to one.** i.e. Given $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$, want $\hat{y}^{(i)} \approx y^{(i)}$.

Simplify $w^T X + b$ to $w^T X$ by adding $x_0 = 1$ and b to be the first element of w .

Loss(error) function: $L(\hat{y}, y)$. -> defined to single training example.

$$L(\hat{y}, y) = -(y \log \hat{y} + (1 - y) \log(1 - \hat{y})).$$

Cost function $J(w, b)$ -> defined to the entity of training examples.

$$J(w, b) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)}).$$

The difference between the cost function and the loss function for logistic regression? -> The loss function computes the error for a single training example

while the cost function is the average of the loss functions of the entire training set.

Gradient Descent

$$J(w, b) = -\frac{1}{m} \sum_{i=1}^m (y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})).$$

Want to find w, b that minimize $J(w, b)$.

After initializing, iterating:

$$\text{Repeat} \{w := w - \alpha \frac{dJ(w, b)}{dw}; b := b - \alpha \frac{dJ(w, b)}{db}\} \text{ until they converge.}$$

Computation Graph: One step of backward propagation on a computation graph yields derivative of final output variable.

An example of 2-dimensional weight vector: (single training example)

$$\{x_1, w_1, x_2, w_2, b\} \rightarrow z = w_1 x_1 + w_2 x_2 + b \iff a = \sigma(z) \iff L(a, y)$$

$$da = \frac{dL(a, y)}{da} = \frac{d(-(y \log a + (1 - y) \log(1 - a)))}{da} = -\frac{y}{a} + \frac{1 - y}{1 - a}$$

$$dz = \frac{dL}{dz} = \frac{dL da}{da dz} = \left(-\frac{y}{a} + \frac{1 - y}{1 - a}\right)(a(1 - a)) = a - y$$

$$dw_1 = \frac{dL}{dw_1} = x_1 dz$$

$$dw_2 = \frac{dL}{dw_2} = x_2 dz$$

$$db = \frac{dL}{db} = dz$$

Update:

$$w_1 := w_1 - \alpha dw_1; w_2 := w_2 - \alpha dw_2; b := b - \alpha db.$$

One single step of Gradient Descent PseudoCode:

```
// Given:
// x_1, x_2, ..., x_m as m training examples, each of them has n dimensions
// called by x_i[j].
// alpha: learning rate
// y_1, y_2, ..., y_m as m labels, each of them is a real number.

J = 0, dw_1 = 0, dw_2 = 0, ..., dw_n = 0, db = 0
initialize w and b // w is as [w_1, w_2, ..., w_n]

for i = 1 to m
    z_i = w.T * x_i + b
```

```

a_i = sigmoid(z_i)
J += -(y_i * log(a_i) + (1 - y_i)*log(1 - a_i))
dz_i = a_i - y_i
dw_1 += x_i[1] * dz_i
dw_2 += x_i[2] * dz_i
...
dw_n += x_i[n] * dz_i
db += dz_i
// average
J = J / m
dw_1 = dw_1 / m
dw_2 = dw_2 / m
...
dw_n = dw_n / m
db = db / m
// update
w_1 = w_1 - alpha * dw_1
w_2 = w_2 - alpha * dw_2
...
w_n = w_n - alpha * dw_n
b = b - alpha * db

```

Weakness: too many for-loops, which cause inefficiency. -> Use vectorization.

vectorization

The art of getting rid of explicit folders in the code. Why the explicit for-loop can be so slow?

```
np.dot(w, x) + b
```

GPU and CPU both have SIMD - “single instruction multiple data”.

Modify the code snippet to the vectorized version.

```

J = 0
dw = np.zeros((n_x, 1))
db = 0

for i = 1 to m
    z_i = w.T * x_i + b
    a_i = sigmoid(z_i)
    J += -(y_i * log(a_i) + (1 - y_i)*log(1 - a_i))
    dz_i = a_i - y_i
    dw += np.dot(x_i, dz_i)
    db += dz_i
# ...

```

Or even better, getting rid of the outer for-loop as well.

$$[z^{(1)}, z^{(2)}, \dots, z^{(m)}] = w^T X + [b, b, \dots, b] = [w^T x^{(1)} + b, w^T x^{(2)} + b, \dots, w^T x^{(m)} + b]$$

```
z = np.dot(w.T, x) + b
a = sigmoid(z)
```

The dimension of X : (n_x, m) .

vectorizing Logistic Regression

$$dz^{(1)} = a^{(1)} - y^{(1)}$$

$$dz^{(2)} = a^{(2)} - y^{(2)}$$

...

$$dZ = [dz^{(1)}, dz^{(2)}, \dots, dz^{(m)}]$$

$$A = [a^{(1)}, a^{(2)}, \dots, a^{(m)}]$$

$$Y = [y^{(1)}, y^{(2)}, \dots, y^{(m)}]$$

$$dZ = A - Y$$

$$db = \frac{1}{m} \sum_{i=1}^m dz^{(i)} = \frac{1}{m} np.sum(dZ)$$

$$dw = \frac{1}{m} X dZ^T \text{ (} n \times 1 \text{ matrix)}.$$

```
Z = np.dot(w.T, x) + b
```

```
A = sigmoid(Z)
```

```
dZ = A - Y
```

```
dw = 1/m * np.dot(X, dZ.T)
```

```
db = 1/m * np.sum(dZ)
```

```
w = w - alpha * dw
```

```
b = b - alpha * db
```